

CS148: Introduction to Computer Graphics and Imaging

Programmable Graphics Pipelines



Topics

The fixed-function graphics pipeline

Programmable stages


- Vertex shaders
- Fragment shaders

GL shading language (GLSL)

Mapping other applications to GPUs

The Graphics Pipeline

A Trip Down the Graphics Pipeline



Application

Simulation

Input event handlers

Modify data structures

Database traversal

Primitive generation

Graphics library utility functions (glu*)

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

Command


Command queue

Command interpretation

Unpack and perform format conversion

Maintain graphics state

```
glLoadIdentity( );
glMultMatrix( T );
glBegin( GL_TRIANGLE_STRIP );
glColor3f( 0.0, 0.5, 0.0 );
glVertex3f( 0.0, 0.0, 0.0 );
glColor3f( 0.5, 0.0, 0.0 );
glVertex3f( 1.0, 0.0, 0.0 );
glColor3f( 0.0, 0.5, 0.0 );
glVertex3f( 0.0, 1.0, 0.0 );
glColor3f( 0.5, 0.0, 0.0 );
glVertex3f( 1.0, 1.0, 0.0 );
...
glEnd( );
```



CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

Vertex (per-vertex)


Vertex transformation

Normal transformation

Texture coordinate generation

Texture coordinate transformation

Lighting (light sources and surface reflection)



CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

Primitive Assembly

Combine transformed/lit vertices into primitives

- 1 vert -> point
- 2 verts -> line
- 3 verts -> triangle

Clipping

Perspective projection

Transform to window coordinates (viewport)

Determine orientation (CW/CCW)

Back-face cull

CS148 Lecture 11




Imran Haque/Pat Hanrahan, Fall 2010

Rasterization

Setup (per-triangle)

Sampling (triangle = {fragments})

Interpolation (interpolate colors and coordinates)



Screen-space triangles

Fragments

CS148 Lecture 11



Imran Haque/Pat Hanrahan, Fall 2010

Texture

Textures are arrays indexed by floats (Sampler)

Texture address calculation

Texture interpolation and filtering



Fragments


Texture Fragments

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

Fragment

Combine texture sampler outputs
Per-fragment shading




CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

Framebuffer Operations

Owner, scissor, depth, alpha and stencil tests
Blending or compositing




CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010



Display

Gamma correction

Analog to digital conversion



Framebuffer Pixels



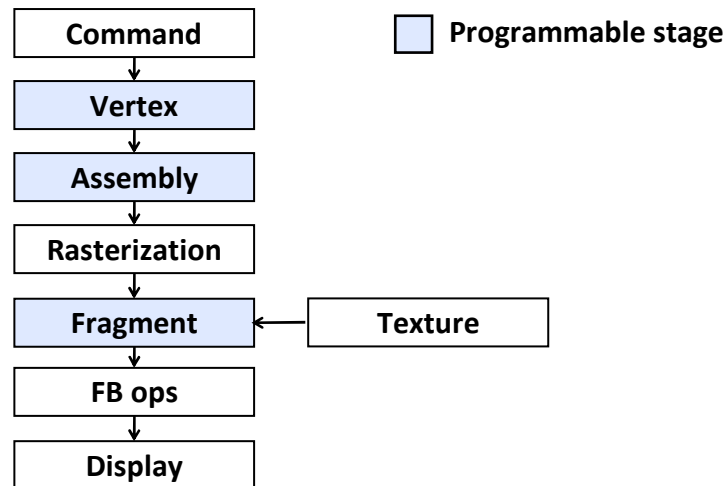
Light

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

Programming Stages


Programmable Graphics Pipeline



CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

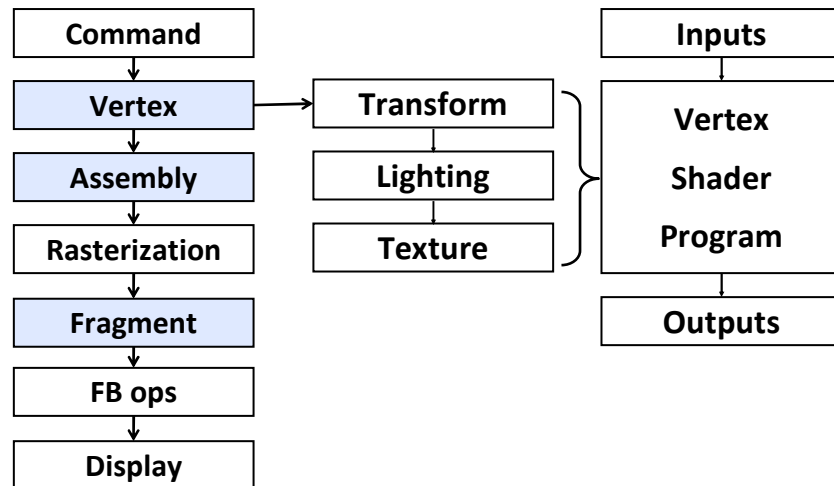
Programmable Graphics Pipeline



CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010


Programmable Graphics Pipeline



CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

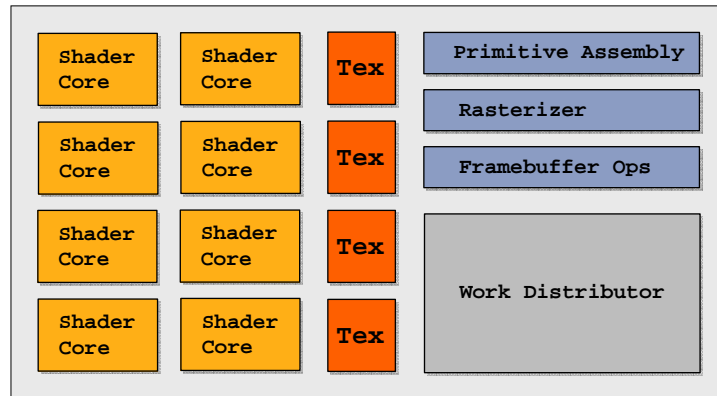
Shader Program Architecture



CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

What's in a GPU?



CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

What's in a GPU?



NVIDIA GF100
(GeForce GTX 480)

AMD Cypress
(Radeon HD 5870)

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

GLSL

Simple Vertex and Fragment Shaders

```
// simple.vert
void main()
{
    gl_Position =
        gl_ModelViewMatrix *
        gl_ProjectionMatrix * gl_Vertex;
    gl_Normal = gl_NormalMatrix * gl_Normal;
    gl_FrontColor = gl_Color;
    gl_BackColor = gl_Color;
}

// simple.frag
void main()
{
    gl_FragColor = gl_Color
}
```

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

Uniform Variables

Uniforms are variables set by the program that can be changed at runtime, but are constant across each execution of the shader;

Changed at most once per primitive

```
// Predefined OpenGL state
uniform mat4 gl_ModelViewMatrix;
uniform mat4 gl_ProjectionMatrix;
uniform mat4 gl_NormalMatrix;

// User-defined
uniform float time;
```

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

Attribute Variables

Attributes variables are properties of a vertex

They are the inputs of the vertex shader

```
attribute vec4 gl_Color;
varying vec4 gl_FrontColor;
varying vec4 gl_BackColor;

void main() {
    gl_FrontColor = gl_Color;
}
```

N. B. All `glVertex* ()` calls result in a `vec4`

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

Varying Variables

Varying variables are the outputs of the vertex shader

```
attribute vec4 gl_Color;  
varying vec4 gl_FrontColor;  
varying vec4 gl_BackColor;  
  
void main() {  
    gl_FrontColor = gl_Color;  
}
```

Varying Variables

The varying variables are interpolated across the triangle

gl_Color is set to gl_FrontColor or gl_BackColor depending on whether the triangle is front facing or back facing

```
varying vec4 gl_Color;  
vec4 gl_FragColor;  
  
void main() {  
    gl_FragColor = gl_Color;  
}
```

Vectors

Constructors

```
vec3 V3 = vec3(1.0, 2.0, 3.0);  
vec4 V4 = vec4(V3, 4.0);
```

Swizzling

```
vec2 V2 = V4.xy;  
vec4 V4Reverse = V4.wzyx;  
vec4 Result = V4.xyzw + V4.xxxx;
```

Basic Vector Operators

```
float Result = dot(V4, V4Reverse);  
vec3 Result = cross(V3, vec3(1.0, 0.0, 0.0));
```

N. B. Points, vectors, normals and colors are all vec's

Textures

```
uniform sampler2D SomeTexture;  
  
void main()  
{  
    vec4 SomeTextureColor =  
        texture2D(SomeTexture, vec2(0.5, 0.5));  
}
```

N. B. Textures coordinates are from (0, 0) to (1, 1)

Communicating with GLSL

Graphics state is available as uniform variables

```
uniform mat4 gl_ModelViewMatrix;
```

Can extend state

```
uniform float x;  
  
addr = GetUniformLocation( program, "x");  
glUniform1f( addr, value );
```

Primitive attributes are available as attribute variables

Can extend attributes (inside glBegin/glEnd)

```
uniform float y;  
  
addr = GetAttributeLocation( program, "y");  
glVertexAttrib1f( addr, value );
```

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

The OpenGL Pipeline in GLSL - Vertex

Built-in attributes

<code>vec4 gl_Vertex</code>	<code>glVertex*()</code>
<code>vec4 gl_Color</code>	<code>glColor*()</code>
<code>vec4 gl_SecondaryColor</code>	<code>glSecondaryColor*()</code>
<code>vec4 gl_Normal</code>	<code>glNormal()</code>
<code>vec4 gl_MultiTexCoord0</code>	<code>glMultiTexCoord(0, ...)</code>

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

The OpenGL Pipeline in GLSL - Fragment

Built-in varying

```
vec4 gl_Position
vec4 gl_FrontColor, gl_BackColor
vec4 gl_FrontSecondaryColor, gl_BackSecondaryColor
vec4 gl_TexCoord[n]
vec4 gl_FragCoord
```

Outputs

```
vec4 gl_FragColor
vec4 gl_FragDepth
```

Simple Pixel Shader

```
varying vec2 TexCoord0;
varying vec2 TexCoord1;
uniform sampler2D SomeTexture0;
uniform sampler2D SomeTexture1;
void main()
{
    gl_FragColor =
        texture2D(SomeTexture0, TexCoord0) * 0.5 +
        texture2D(SomeTexture1, TexCoord1) * 0.5;
}
```

This makes it easy to build image processing filters

Limitations

Memory


- No access to neighboring fragments
- Limited stack space, instruction count
- Cannot read and write framebuffer

Performance

- Branching support is limited and slow
- Graphics card will timeout if code takes too long
- Variable support across different graphics cards

GPU Computing

Why GPGPU?



GPU's are great if problem:

- Executes the same code many times on different input
- Needs lots of math
- Does not share data between executing components
- Has lots of work to do without CPU intervention

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

Computation on GPU's

Beyond basic graphics pipeline

- Collision detection
- Fluid and cloth simulation
- Physics
- Ray-tracing

Beyond graphics


- Protein folding (Folding@Home)
- Speech recognition
- Partial differential equation solvers
- Fourier transforms

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

An Example GPGPU Application - PAPER

- **Molecular overlay optimization:** used in computational drug discovery to find new active compounds from a database given one active “query” molecule

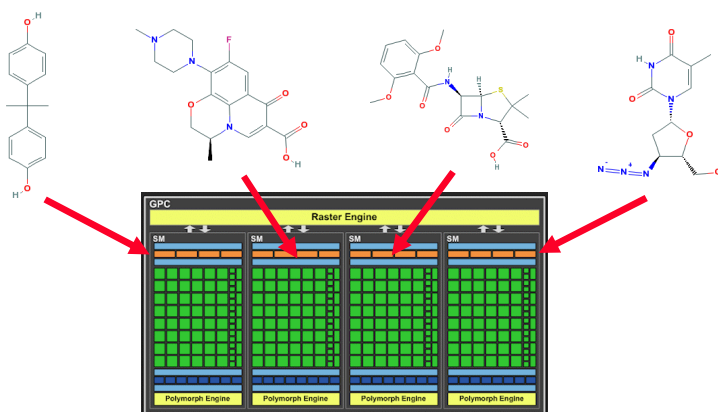


- Complexity $O(MN)$: double-loop over all atom pairs
- DB = ~10M molecules; CPU = 10ms/overlay = ~2 days/query
- *Use GPU to exploit parallelism of problem.*

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

GPU Parallelism 1




Each optimization is independent, and each SM (OpenCL work-group) executes independently, so run one DB molecule per GPU core

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

GPU Parallelism 2

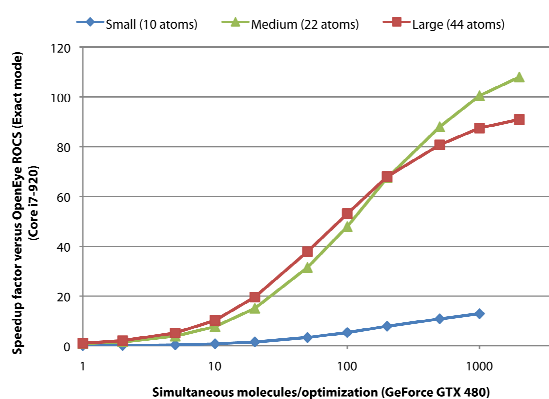


GPU cores have wide internal parallelism. Each atom pair in an optimization is independent – map each to a shader unit (OpenCL work-item), and loop.

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010

GPGPU Conclusion



>100x speedup if there's lots of parallel work

48 hr for CPU DB search -> 30-60 min with GPU!

CS148 Lecture 11

Imran Haque/Pat Hanrahan, Fall 2010